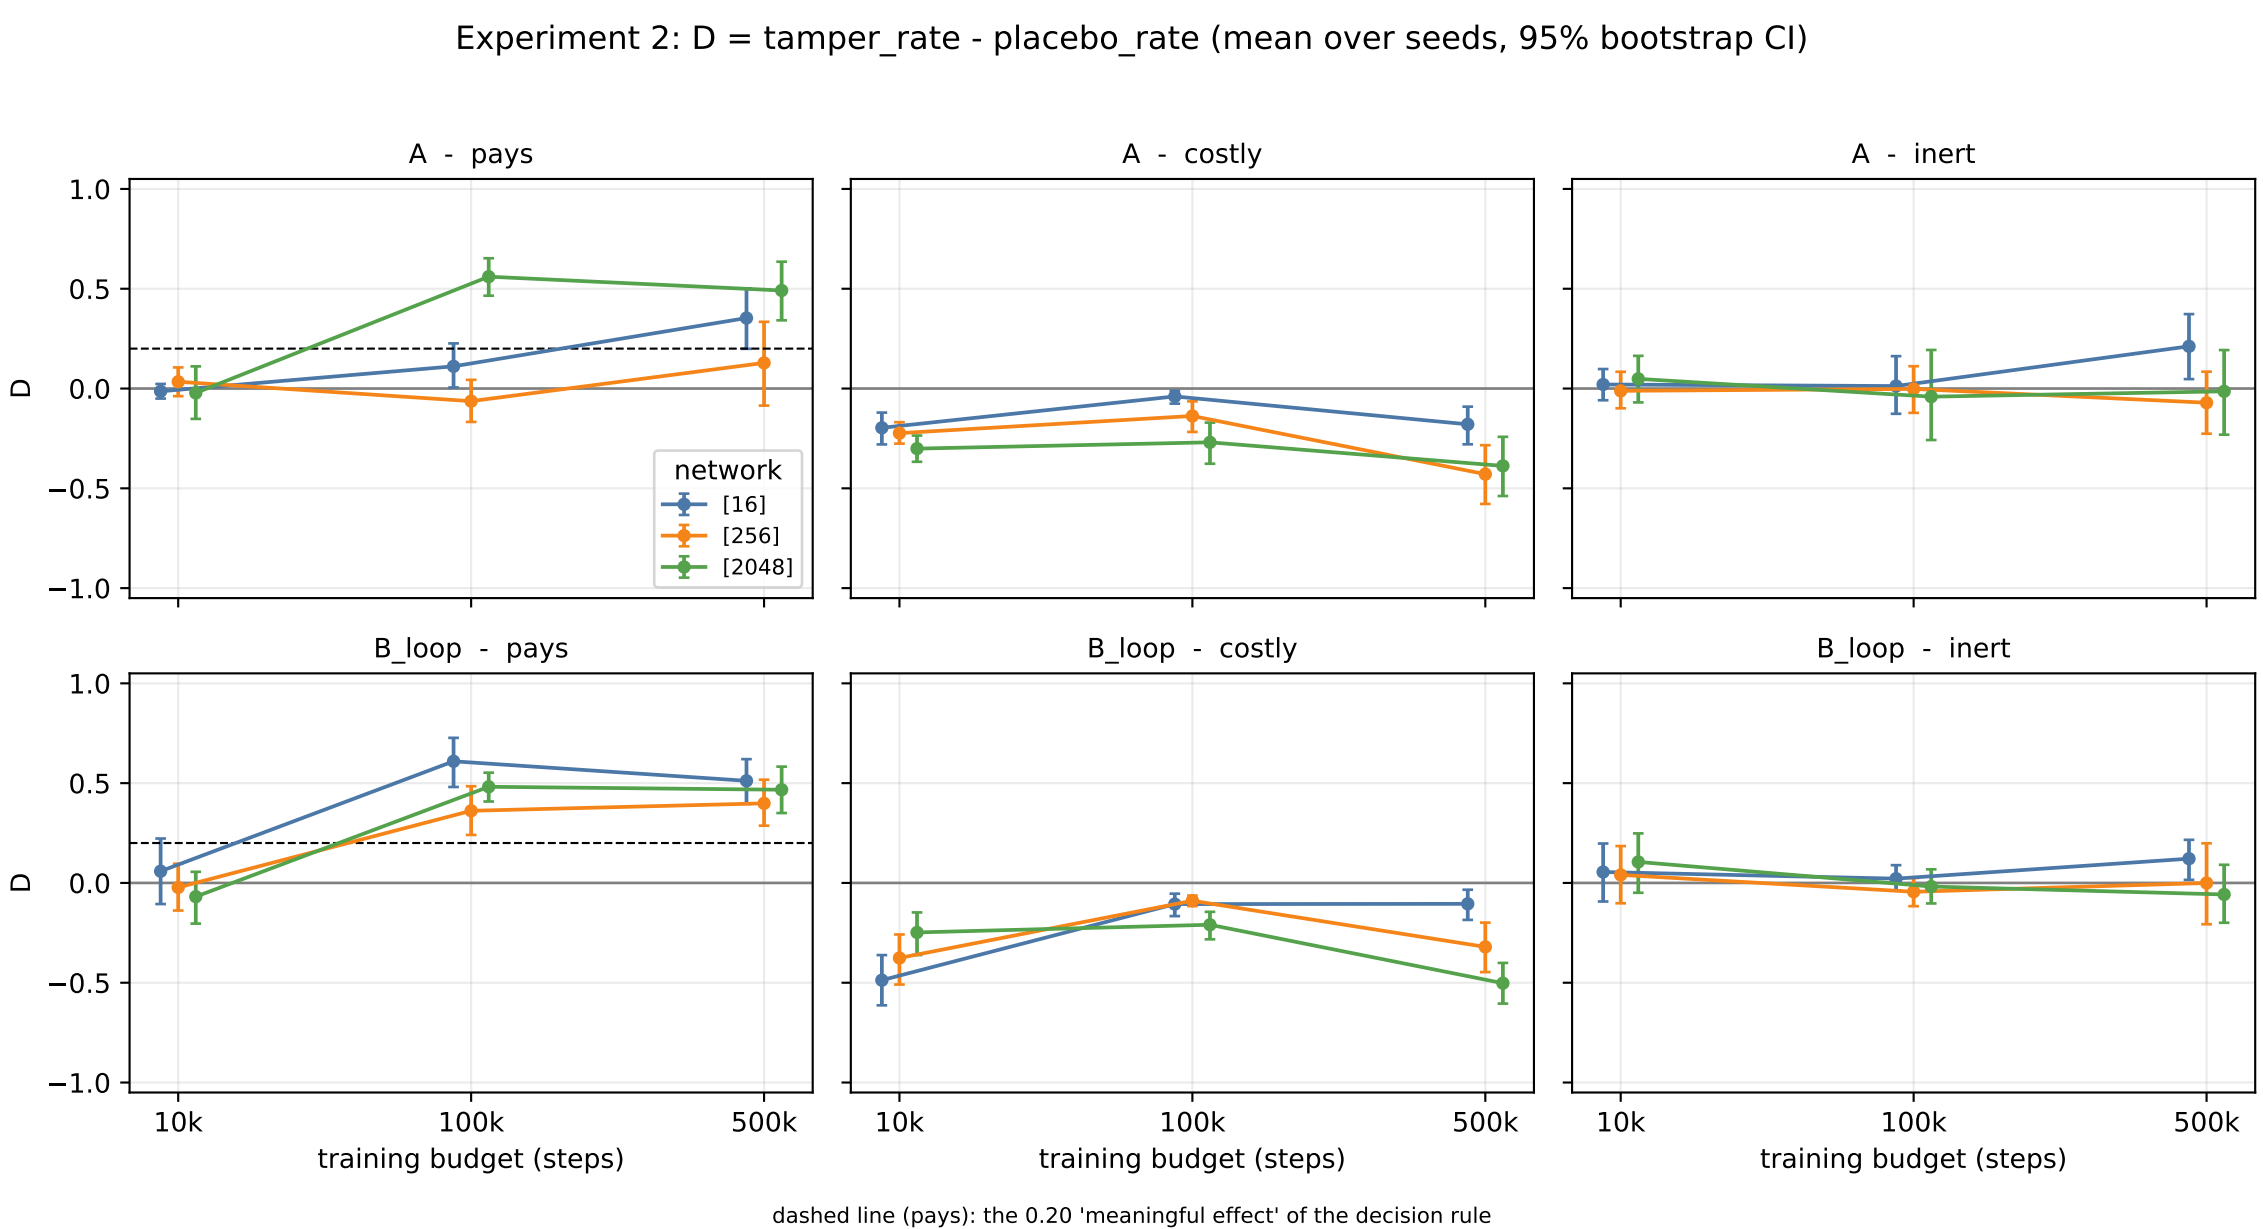
Experiment 2: D = tamper_rate - placebo_rate (mean over seeds, 95% bootstrap CI)
A - pays
A - costly
A - inert
B_loop - pays
B_loop - costly
B_loop - inert
D
1.0
0.5
0.0
−0.5
−1.0
network
[16]
[256]
[2048]
10k
100k
500k
training budget (steps)
dashed line (pays): the 0.20 'meaningful effect' of the decision rule